

```
In [230]: import Pkg;
#Pkg.add("CSV")
#Pkg.add("Random")
#Pkg.add("PyPlot")
#Pkg.add("MLLabelUtils")
#Pkg.add("MLDataUtils")
#Pkg.add("DecisionTree")
#Pkg.add("ScikitLearn")
#Pkg.add("XGBoost")
#Pkg.add("PyCall")
#Pkg.add("MLBase")
#Pkg.add("LearnBase")
```

Diagnosing Breast Cancer using Julia

Loading necessary packages and reading the [dataset \(https://www.kaggle.com/uciml/breast-cancer-wisconsin-data\)](https://www.kaggle.com/uciml/breast-cancer-wisconsin-data).

```
In [231]: using CSV
using DataFrames
df = CSV.read("wdbc.csv");
```

Setting Random Seed to produce same output everytime, it helps in debugging

```
In [232]: RANDOM_STATE = 42
using Random
Random.seed!(RANDOM_STATE)
```

Out[232]: MersenneTwister(UInt32[0x0000002a], Random.DSFMT.DSFMT_state(Int32[964434469, 1073036706, 1860149520, 1073503458, 1687169063, 1073083486, -399267803, 1072983952, -909620556, 1072836235 ... -293054293, 1073002412, -1300127419, 1073642642, 1917177374, -666058738, -337596527, 1830741494, 382, 0]), [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0 ... 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0], UInt128[0x00000000000000000000000000000000, 0x00000000000000000000000000000000], 1002, 0)

First 5 samples of the dataset

```
In [233]: first(df,5)
```

Out[233]: 5 rows × 32 columns (omitted printing of 25 columns)

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
	Int64	String	Float64	Float64	Float64	Float64	Float64
1	842302	M	17.99	10.38	122.8	1001.0	0.1184
2	842517	M	20.57	17.77	132.9	1326.0	0.08474
3	84300903	M	19.69	21.25	130.0	1203.0	0.1096
4	84348301	M	11.42	20.38	77.58	386.1	0.1425
5	84358402	M	20.29	14.34	135.1	1297.0	0.1003

Printing the number of instances and number of columns

```
In [234]: # size function returns the number of row and columns of the dataset
row, col = size(df)
println("Number of Instances: ", row)
println("Number of Columns: ", col)
```

Number of Instances: 569
Number of Columns: 32

Headers / Column names of the dataset

In [235]:

println(names(df))

Symbol[:id, :diagnosis, :radius_mean, :texture_mean, :perimeter_mean, :area_mean, :smoothness_mean, :compactness_mean, :concavity_mean, Symbol("concave points_mean"), :symmetry_mean, :fractal_dimension_mean, :radius_se, :texture_se, :perimeter_se, :area_se, :smoothness_se, :compactness_se, :concavity_se, Symbol("concave points_se"), :symmetry_se, :fractal_dimension_se, :radius_worst, :texture_worst, :perimeter_worst, :area_worst, :smoothness_worst, :compactness_worst, :concavity_worst, Symbol("concave points_worst"), :symmetry_worst, :fractal_dimension_worst]

Description of dataset

In [236]:

println(describe(df))

32x8 DataFrame									
	Row	variable	mean	min	median	max	nunique	nmissing	eltype
type		Symbol	Union{...}	Any	Union{...}	Any	Union{...}	Nothing	DataT
	1	id	3.03718e7	8670	906024.0	911320502			Int64
	2	diagnosis		B		M	2		String
	3	radius_mean	14.1273	6.981	13.37	28.11			Float64
	4	texture_mean	19.2896	9.71	18.84	39.28			Float64
	5	perimeter_mean	91.969	43.79	86.24	188.5			Float64
	6	area_mean	654.889	143.5	551.1	2501.0			Float64
	7	smoothness_mean	0.0963603	0.05263	0.09587	0.1634			Float64
	8	compactness_mean	0.104341	0.01938	0.09263	0.3454			Float64
	9	concavity_mean	0.0887993	0.0	0.06154	0.4268			Float64
	10	concave points_mean	0.0489191	0.0	0.0335	0.2012			Float64
	11	symmetry_mean	0.181162	0.106	0.1792	0.304			Float64
	12	fractal_dimension_mean	0.0627976	0.04996	0.06154	0.09744			Float64
	13	radius_se	0.405172	0.1115	0.3242	2.873			Float64
	14	texture_se	1.21685	0.3602	1.108	4.885			Float64
	15	perimeter_se	2.86606	0.757	2.287	21.98			Float64
	16	area_se	40.3371	6.802	24.53	542.2			Float64
	17	smoothness_se	0.00704098	0.001713	0.00638	0.03113			Float64
	18	compactness_se	0.0254781	0.002252	0.02045	0.1354			Float64
	19	concavity_se	0.0318937	0.0	0.02589	0.396			Float64
	20	concave points_se	0.0117961	0.0	0.01093	0.05279			Float64
	21	symmetry_se	0.0205423	0.007882	0.01873	0.07895			Float64
	22	fractal_dimension_se	0.0037949	0.0008948	0.003187	0.02984			Float64
	23	radius_worst	16.2692	7.93	14.97	36.04			Float64
	24	texture_worst	25.6772	12.02	25.41	49.54			Float64
	25	perimeter_worst	107.261	50.41	97.66	251.2			Float64
	26	area_worst	880.583	185.2	686.5	4254.0			Float64
	27	smoothness_worst	0.132369	0.07117	0.1313	0.2226			Float64
	28	compactness_worst	0.254265	0.02729	0.2119	1.058			Float64
	29	concavity_worst	0.272188	0.0	0.2267	1.252			Float64
	30	concave points_worst	0.114606	0.0	0.09993	0.291			Float64
	31	symmetry_worst	0.290076	0.1565	0.2822	0.6638			Float64
	32	fractal_dimension_worst	0.0839458	0.05504	0.08004	0.2075			Float64


Frequency of each class

```
In [237]: using FreqTables
frequency_digonosis = freqtable(df[:diagnosis])
```

```
Out[237]: 2-element Named Array{Int64,1}
Dim1
-----
B      357
M      212
```

Histogram of the frequency of each class

```
In [238]: using PyPlot
PyPlot.plt.hist(df[:diagnosis], 2)
plt.title("Histogram of Diagnosis")
plt.xlabel("Malignant(M), Benign(B)")
plt.ylabel("Frequency")
```



```
Out[238]: PyObject Text(24.000000000000007, 0.5, 'Frequency')
```

Data Preprocessing

Converting the categorical feature "diagnosis" to integer by doing label encoding

```
In [239]: using MLLabelUtils
y = convertlabel(LabelEnc.MarginBased,df[:diagnosis])
y = classify.(y, LabelEnc.ZeroOne(Int,1))
newdf = copy(df)
deletecols!(newdf, :diagnosis)
newdf[:diagnosis] = y;
```

```
In [240]: #deletecols!(newdf, :id)
```

Dataset after encoding

```
In [241]: first(newdf,5)
```

```
Out[241]: 5 rows × 32 columns (omitted printing of 26 columns)
```

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
	Int64	Float64	Float64	Float64	Float64	Float64
1	842302	17.99	10.38	122.8	1001.0	0.1184
2	842517	20.57	17.77	132.9	1326.0	0.08474
3	84300903	19.69	21.25	130.0	1203.0	0.1096
4	84348301	11.42	20.38	77.58	386.1	0.1425
5	84358402	20.29	14.34	135.1	1297.0	0.1003

Description of Dataset after encoding

In [242]: `println(describe(newdf))`

32×8 DataFrame									
	Row	variable	mean	min	median	max	nunique	nmissing	eltype
		Symbol	Float64	Real	Float64	Real	Nothing	Nothing	DataT
type									
	1	id	3.03718e7	8670	906024.0	911320502			Int64
	2	radius_mean	14.1273	6.981	13.37	28.11			Float
64	3	texture_mean	19.2896	9.71	18.84	39.28			Float
64	4	perimeter_mean	91.969	43.79	86.24	188.5			Float
64	5	area_mean	654.889	143.5	551.1	2501.0			Float
64	6	smoothness_mean	0.0963603	0.05263	0.09587	0.1634			Float
64	7	compactness_mean	0.104341	0.01938	0.09263	0.3454			Float
64	8	concavity_mean	0.0887993	0.0	0.06154	0.4268			Float
64	9	concave points_mean	0.0489191	0.0	0.0335	0.2012			Float
64	10	symmetry_mean	0.181162	0.106	0.1792	0.304			Float
64	11	fractal_dimension_mean	0.0627976	0.04996	0.06154	0.09744			Float
64	12	radius_se	0.405172	0.1115	0.3242	2.873			Float
64	13	texture_se	1.21685	0.3602	1.108	4.885			Float
64	14	perimeter_se	2.86606	0.757	2.287	21.98			Float
64	15	area_se	40.3371	6.802	24.53	542.2			Float
64	16	smoothness_se	0.00704098	0.001713	0.00638	0.03113			Float
64	17	compactness_se	0.0254781	0.002252	0.02045	0.1354			Float
64	18	concavity_se	0.0318937	0.0	0.02589	0.396			Float
64	19	concave points_se	0.0117961	0.0	0.01093	0.05279			Float
64	20	symmetry_se	0.0205423	0.007882	0.01873	0.07895			Float
64	21	fractal_dimension_se	0.0037949	0.0008948	0.003187	0.02984			Float
64	22	radius_worst	16.2692	7.93	14.97	36.04			Float
64	23	texture_worst	25.6772	12.02	25.41	49.54			Float
64	24	perimeter_worst	107.261	50.41	97.66	251.2			Float
64	25	area_worst	880.583	185.2	686.5	4254.0			Float
64	26	smoothness_worst	0.132369	0.07117	0.1313	0.2226			Float
64	27	compactness_worst	0.254265	0.02729	0.2119	1.058			Float
64	28	concavity_worst	0.272188	0.0	0.2267	1.252			Float
64	29	concave points_worst	0.114606	0.0	0.09993	0.291			Float
64	30	symmetry_worst	0.290076	0.1565	0.2822	0.6638			Float
64	31	fractal_dimension_worst	0.0839458	0.05504	0.08004	0.2075			Float
64	32	diagnosis	0.372583	0	0.0	1			Int64

Splitting the dataset into train, test

Shuffling the Dataset

```
In [243]: using MLDataUtils
newdf_s = shuffleobs(newdf);
first(newdf_s,10)
```

Out[243]: 10 rows × 32 columns (omitted printing of 26 columns)

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
	Int64	Float64	Float64	Float64	Float64	Float64
1	914333	14.87	20.21	96.12	680.9	0.09587
2	864292	10.51	20.19	68.64	334.2	0.1122
3	903554	12.1	17.72	78.07	446.2	0.1029
4	894604	10.25	16.18	66.52	324.2	0.1061
5	86517	18.66	17.12	121.4	1077.0	0.1054
6	894329	9.042	18.9	60.07	244.5	0.09968
7	89382601	14.61	15.69	92.68	664.9	0.07618
8	891923	13.77	13.27	88.06	582.7	0.09198
9	915452	16.3	15.7	104.7	819.8	0.09427
10	917897	9.847	15.68	63.0	293.2	0.09492

Spliting into train and test

```
In [244]: train, test = splitobs(newdf_s, at = 0.70);
```

Seperating the features and class label

```
In [245]: row,col = size(train)
```

Out[245]: (398, 32)

```
In [246]: X_train = convert(Matrix, train[:,2:col-1]);
y_train = convert(Matrix, train[:,col:col]);
X_test = convert(Matrix, test[:,2:col-1]);
y_test = convert(Matrix, test[:,col:col]);
```

Converting (n,1) dimension to (n,) as per the classifiers requirements

```
In [247]: y_train = vec(y_train);
y_test = vec(y_test);
```

Applying algorithms to train models

RandomForest

Loading necessary packages

```
In [248]: using DecisionTree
```

```
In [249]: ?RandomForestClassifier
```

search: **RandomForestClassifier**

```
Out[249]: RandomForestClassifier(; n_subfeatures::Int=-1,
                                   n_trees::Int=10,
                                   partial_sampling::Float=0.7,
                                   max_depth::Int=-1,
                                   rng=Random.GLOBAL_RNG)
```

Random forest classification. See [DecisionTree.jl's documentation \(https://github.com/bensadeghi/DecisionTree.jl\)](https://github.com/bensadeghi/DecisionTree.jl).

Hyperparameters:

- `n_subfeatures` : number of features to consider at random per split (default: -1, sqrt(# features))
- `n_trees` : number of trees to train (default: 10)
- `partial_sampling` : fraction of samples to train each tree on (default: 0.7)
- `max_depth` : maximum depth of the decision trees (default: no maximum)
- `min_samples_leaf` : the minimum number of samples each leaf needs to have
- `min_samples_split` : the minimum number of samples in needed for a split
- `min_purity_increase` : minimum purity needed for a split
- `rng` : the random number generator to use. Can be an `Int` , which will be used to seed and create a new random number generator.

Implements `fit!` , `predict` , `predict_proba` , `get_classes`

Creating an object of Classifier

```
In [250]: using DecisionTree
using ScikitLearn: fit!, predict
rfc = RandomForestClassifier(n_trees = 110, n_subfeatures = 15, max_depth = 7);
```

Fitting the classifiers on training data samples

```
In [251]: fit!(rfc, X_train, y_train);
```

predicting the test data samples

```
In [252]: y_pred_rfc = predict(rfc, X_test);
println("A portion of prediction:")
y_pred_rfc[1:10]
```

A portion of prediction:

```
Out[252]: 10-element Array{Int64,1}:
 0
 0
 0
 0
 0
 1
 1
 0
 0
 1
```

```
In [253]: cm = confusion_matrix(y_test, y_pred_rfc)

2×2 Array{Int64,2}:
 96  2
 7  66
```

```
Out[253]: Classes:  [0, 1]
Matrix:
Accuracy: 0.9473684210526315
Kappa:    0.8914898117464569
```

AdaBoostClassifier

```
In [254]: ?AdaBoostStumpClassifier
```

search: **AdaBoostStumpClassifier**

```
Out[254]: AdaBoostStumpClassifier(; n_iterations::Int=0)
```

Adaboosted decision tree stumps. See [DecisionTree.jl's documentation \(https://github.com/bensadeghi/DecisionTree.jl\)](https://github.com/bensadeghi/DecisionTree.jl).

Hyperparameters:

- `n_iterations` : number of iterations of AdaBoost
- `rng` : the random number generator to use. Can be an `Int` , which will be used to seed and create a new random number generator.

Implements `fit!` , `predict` , `predict_proba` , `get_classes`

Parameter Tuning

```
In [255]: using ScikitLearn.GridSearch: GridSearchCV

gridsearch = GridSearchCV(AdaBoostStumpClassifier(), Dict{:n_iterations => 1:1:100})
fit!(gridsearch, X_train, y_train)
println("Best parameters: $(gridsearch.best_params_)" )
```

Best parameters: Dict{Symbol,Any}{:n_iterations=>65}

Creating object of Random Forest

```
In [256]: abc = AdaBoostStumpClassifier(n_iterations = 89);
fit!(abc, X_train, y_train);
```

```
In [257]: y_pred_abc = predict(abc, X_test);
println("A portion of prediction:")
y_pred_abc[1:10]
```

A portion of prediction:

```
Out[257]: 10-element Array{Int64,1}:
 0
 0
 0
 0
 0
 0
 0
 1
 0
 1
 1
```

```
In [258]: cm = confusion_matrix(y_test, y_pred_abc)
```

2x2 Array{Int64,2}:
97 1
6 67

```
Out[258]: Classes: [0, 1]
Matrix:
Accuracy: 0.9590643274853801
Kappa: 0.915603186913911
```

```
In [259]: ?DecisionTreeClassifier
```

search: **DecisionTreeClassifier**

```
Out[259]: DecisionTreeClassifier(; pruning_purity_threshold=0.0,
                                max_depth::Int=-1,
                                min_samples_leaf::Int=1,
                                min_samples_split::Int=2,
                                min_purity_increase::Float=0.0,
                                n_subfeatures::Int=0,
                                rng=Random.GLOBAL_RNG)
```

Decision tree classifier. See [DecisionTree.jl's documentation \(https://github.com/bensadeghi/DecisionTree.jl\)](https://github.com/bensadeghi/DecisionTree.jl)

Hyperparameters:

- `pruning_purity_threshold` : (post-pruning) merge leaves having `>=thresh` combined purity (default: no pruning)
- `max_depth` : maximum depth of the decision tree (default: no maximum)
- `min_samples_leaf` : the minimum number of samples each leaf needs to have (default: 1)
- `min_samples_split` : the minimum number of samples in needed for a split (default: 2)
- `min_purity_increase` : minimum purity needed for a split (default: 0.0)
- `n_subfeatures` : number of features to select at random (default: keep all)
- `rng` : the random number generator to use. Can be an `Int` , which will be used to seed and create a new random number generator.

Implements `fit!` , `predict` , `predict_proba` , `get_classes`

```
In [260]: dtc = DecisionTreeClassifier()
fit!(dtc, X_train, y_train);
y_pred_dtc = predict(dtc, X_test);
println("A portion of prediction:")
y_pred_dtc[1:10]
```

A portion of prediction:

```
Out[260]: 10-element Array{Int64,1}:
 0
 0
 0
 0
 0
 0
 0
 1
 0
 0
 1
```

```
In [261]: cm = confusion_matrix(y_test, y_pred_dtc)
```

2x2 Array{Int64,2}:
96 2
10 63

```
Out[261]: Classes: [0, 1]
Matrix:
Accuracy: 0.9298245614035088
Kappa: 0.8545506095832152
```

XGBoost

```
In [262]: #import Pkg;
#Pkg.add("XGBoost")
```

```
In [263]: using XGBoost
```

```
In [264]: ?xgboost
```

search: **xgboost XGBoost**

```
Out[264]: No documentation found.

XGBoost.xgboost is a Function .
# 1 method for generic function "xgboost":
[1] xgboost(data, nrounds::Integer; label, param, watchlist, metrics, obj, feval, group, kwargs...) in XGBoost at C:\Users\idpau\.julia\packages\XGBoost\LXjD0\src\xgboost_lib.jl:145
```



```
In [265]: num_round = 5
xgbc = xgboost(X_train, num_round, label = y_train, eta = 1, max_depth = 2)
```

```
[1]      train-rmse:0.180252
[2]      train-rmse:0.147482
[3]      train-rmse:0.129229
[4]      train-rmse:0.123099
[5]      train-rmse:0.116709
```

Out[265]: `Booster(Ptr{Nothing} @0x00000000307d3490)`

```
In [266]: y_pred_xgbc = XGBoost.predict(xgbc, X_test)
print("test-error=", sum((y_pred_xgbc .> 0.5) .!= y_test) / float(size(y_pred_xgbc)[1]), "\n")
```

test-error=0.10526315789473684

```
In [267]: y_pred_xgbc
```

Out[267]: 171-element Array{Float32,1}:
0.0123640895
0.0123640895
0.0123640895
-0.35979003
0.0123640895
0.4310953
0.9815798
0.0123640895
0.0123640895
0.9815798
0.0123640895
0.26188624
1.028301
:
0.4447895
1.028301
0.89318776
0.0123640895
0.0123640895
0.0123640895
0.0123640895
0.0123640895
0.9815798
0.67093486
0.9815798
0.9815798

```
In [268]: cm = confusion_matrix(y_test, y_pred_xgbc[0])
```

BoundsError: attempt to access 171-element Array{Float32,1} at index [0]

Stacktrace:
[1] getindex(::Array{Float32,1}, ::Int64) at .\array.jl:729
[2] top-level scope at In[268]:1

Performance and Results

Performance evaluation criteria

Defining functions for calculating different performance metrics

```
In [269]: using PyCall
math = pyimport("math")

function accuracy_(tn,fp,fn,tp)
    return ((tp+tn)/(tp+fp+fn+tn))
end

#True Positive Rate or Recall
function sensitivity(tp,fn)
    return (tp / (tp + fn))
end

tprate = sensitivity
recall = sensitivity

#True Negative Rate
function specificity(tn,fp)
    return (tn / (fp + tn))
end

function mcc(tn,fp,fn,tp)
    return (tp*tn-fp*fn)/math.sqrt((tp+fp)*(tp+fn)*(tn+fp)*(tn+fn))
end

function auc_score(tn,fp,fn,tp)
    return (sensitivity(tp,fn) + specificity(tn,fp)) / 2
end

function gmean_score(tn,fp,fn,tp)
    return math.sqrt(sensitivity(tp,fn) * specificity(tn,fp))
end

function precision(tp, fp)
    return (tp / (tp + fp))
end

function f1score(tp, fp, fn)
    return (2*tp/(2*tp + fp + fn))
end

function fprate(tn,fp)
    return (fp / (fp + tn))
end
```

Out[269]: fprate (generic function with 1 method)

Defininf a function for producing all scores at once.

```
In [270]: function getAllScore(tp,tn,fp,fn)
    acc = accuracy_(tn,fp,fn,tp)
    sn = sensitivity(tp,fn)
    sp = specificity(tn,fp)
    auc = auc_score(tn,fp,fn,tp)
    gmean = gmean_score(tn,fp,fn,tp)
    preci = precision(tp, fp)
    f1 = f1score(tp,fp,fn)
    fpr = fprate(tn,fp)
    return acc, sn, sp, auc, gmean, preci, f1, fpr
end
```

Out[270]: getAllScore (generic function with 1 method)

Result of RandomForest

Loading MLBase to compute roc performance parameter

```
In [271]: using MLBase
r_rfc = roc(y_test, y_pred_rfc)
```

Out[271]: ROCNums{Int64}
p = 73
n = 98
tp = 66
tn = 96
fp = 2
fn = 7

Calling the getAllScore function

```
In [272]: allScoresTuples = getAllScore(r_rfc.tp, r_rfc.tn, r_rfc.fp, r_rfc.fn)
allScoresArray = [i for i in allScoresTuples]
allScores = transpose(allScoresArray)
```

```
Out[272]: 1x8 LinearAlgebra.Transpose{Float64,Array{Float64,1}}:
 0.947368  0.90411  0.979592  0.941851  ...  0.970588  0.93617  0.0204082
```

```
In [273]: ?CSV.write
```

```
Out[273]: CSV.write(file, table; kwargs...) => file
table |> CSV.write(file; kwargs...) => file
```

Write a [Tables.jl interface input \(https://github.com/JuliaData/Tables.jl\)](https://github.com/JuliaData/Tables.jl) to a csv file, given as an IO argument or String /FilePaths.jl type representing the file name to write to.

Supported keyword arguments include:

- `delim::Union{Char, String}=', '` : a character or string to print out as the file's delimiter
- `quotechar::Char='"'` : ascii character to use for quoting text fields that may contain delimiters or newlines
- `openquotechar::Char` : instead of `quotechar` , use `openquotechar` and `closequotechar` to support different starting and ending quote characters
- `escapechar::Char='\"'` : ascii character used to escape quote characters in a text field
- `missingstring::String=""` : string to print for missing values
- `dateformat=Dates.default_format(T)` : the date format string to use for printing out Date & DateTime columns
- `append=false` : whether to append writing to an existing file/IO, if `true` , it will not write column names by default
- `writeheader=!append` : whether to write an initial row of delimited column names, not written by default if appending
- `header` : pass a list of column names (Symbols or Strings) to use instead of the column names of the input table
- `newline='\n'` : character or string to use to separate rows (lines in the csv file)
- `quotestrings=false` : whether to force all strings to be quoted or not
- `decimal='.'` : character to use as the decimal point when writing floating point numbers

```
In [274]: allScoresArray
```

```
Out[274]: 8-element Array{Float64,1}:
 0.9473684210526315
 0.9041095890410958
 0.9795918367346939
 0.9418507128878948
 0.941094242325505
 0.9705882352941176
 0.9361702127659575
 0.02040816326530612
```

Storing the all scores into a dataframe

```
In [275]: columnNames = ["accuracy", "sensitivity", "specificity", "auc", "gmean", "precision", "f1_score", "fpr"]
resultDF = DataFrame(allScores, Symbol.(columnNames), )
```

```
Out[275]: 1 rows × 8 columns
```

	accuracy	sensitivity	specificity	auc	gmean	precision	f1_score	fpr
	Float64	Float64	Float64	Float64	Float64	Float64	Float64	Float64
1	0.947368	0.90411	0.979592	0.941851	0.941094	0.970588	0.93617	0.0204082

Result of AdaBoost

```
In [276]: using MLBase
r_abc = roc(y_test, y_pred_abc)
```

```
Out[276]: ROCNums{Int64}
 p = 73
 n = 98
 tp = 67
 tn = 97
 fp = 1
 fn = 6
```

Calling the getAllScore function

```
In [277]: allScoresTuples = getAllScore(r_abc.tp, r_abc.tn, r_abc.fp, r_abc.fn)
allScoresArray = [i for i in allScoresTuples]
allScores = transpose(allScoresArray)
```

```
Out[277]: 1x8 LinearAlgebra.Transpose{Float64,Array{Float64,1}}:
 0.959064  0.917808  0.989796  0.953802  ...  0.985294  0.950355  0.0102041
```

In [278]:

allScoresArray

Out[278]:

8-element Array{Float64,1}:
0.9590643274853801
0.9178082191780822
0.9897959183673469
0.9538020687727146
0.9531226726851424
0.9852941176470589
0.950354609929078
0.01020408163265306

Storing the all scores into a dataframe

In [279]:

columnNames = ["accuracy", "sensitivity", "specificity", "auc", "gmean", "precision", "f1_score", "fpr"]
resultDF = DataFrame(allScores, Symbol.(columnNames),)

Out[279]:

1 rows × 8 columns

	accuracy	sensitivity	specificity	auc	gmean	precision	f1_score	fpr
	Float64	Float64	Float64	Float64	Float64	Float64	Float64	Float64
1	0.959064	0.917808	0.989796	0.953802	0.953123	0.985294	0.950355	0.0102041

Drawing a ROC curve with auc score

In [280]:

fprArr = [0, fprate(r_rfc.tn,r_rfc.fp), 1]
tprArr = [0, tprate(r_rfc.tp,r_rfc.fn), 1]
auc_ = round(auc_score(r_rfc.tn,r_rfc.fp,r_rfc.fn,r_rfc.tp), digits=3)
plt.plot(fprArr,tprArr,label= string("RandomForest", " ROC (auc =", auc_, ")"))

fprArr = [0, fprate(r_abc.tn,r_abc.fp), 1]
tprArr = [0, tprate(r_abc.tp,r_abc.fn), 1]
auc_ = round(auc_score(r_abc.tn,r_abc.fp,r_abc.fn,r_abc.tp), digits=3)
plt.plot(fprArr,tprArr,label= string("AdaBoost", " ROC (auc =", auc_, ")"))

xlabel("1 - Specificity or (False Positive Rate)")
ylabel("Sensitivity(True Positive Rate)")
plt.plot([0, 1], [0, 1],"r--")
plt.xlim([0.0, 1.05])
plt.ylim([0.0, 1.05])
PyPlot.title("Receiver Operating Characteristic")
grid("on")
legend()
plt.savefig("roc_curve.png",bbox_inches="tight", format="png", dpi=1200)
plt.show() # Display

